

Explaining “Explaining Away”

Michael P. Wellman and Max Henrion

Abstract—*Explaining away* is a common pattern of reasoning in which the confirmation of one cause of an observed or believed event reduces the need to invoke alternative causes. The opposite of explaining away also can occur, where the confirmation of one cause *increases* belief in another. We provide a general qualitative probabilistic analysis of intercausal reasoning and identify the property of the interaction among the causes (*product synergy*) that determines which form of reasoning is appropriate. Product synergy extends the qualitative probabilistic network (QPN) formalism to support qualitative intercausal inference about the directions of change in probabilistic belief. The intercausal relation also justifies Occam's razor, facilitating pruning in the search for likely diagnoses.

I. EXPLAINING AWAY

Keeping track of the dependency or causal structure among events is critical in uncertain reasoning. One fundamental reason is the inherent asymmetry between *predictive* (or causal) reasoning, from cause to effect, and *diagnostic* (or evidential) reasoning, from effect to cause. Pearl [9] clearly illustrates this asymmetry with the “sprinkler” example, which is depicted in Fig. 1. Either *A*, “it rained last night,”¹ or *B*, “the sprinkler was on last night,” could cause *C*, “the grass is wet.” *C* could in turn cause *E*, “the grass is cold and shiny,” as well as *F*, “my shoes are wet.”

Observation of one effect *E*, cold and shiny grass, is evidence for *C*, wet grass, and predicts the other effect *F*, wet shoes. Confirmation of one cause *A*, rain, also leads to the expectation of *C*, wet grass. However, it does *not* provide any evidence for the alternate cause *B*, sprinkling. Suppose prior observation of wet grass had led to defeasible acceptance of sprinkling. In a default reasoning scheme, confirmation of rain should lead to a *retraction* of the hypothesis that the sprinkler had been on. In a probabilistic reasoning scheme, it should lead to a *reduced probability* of the sprinkler hypothesis, even though the possibility of simultaneous sprinkling and rain is allowed.

This common and intuitively compelling pattern of reasoning is called *explaining away* because one cause explains the observed effect and, therefore, reduces the need to invoke other causes. This qualitative pattern of reasoning is entirely compatible with Bayesian inference when probabilistic influences reflect causal relationships [3], [9]. It is also the essence of Occam's razor: slice away hypotheses that are unnecessary to account for the evidence. Indeed, Paek [8] applies minimization of causal justifications to realize the explaining away pattern in a circumscriptive logic.

Pearl [9] uses the revealed asymmetry of inference with respect to causal direction to argue for incorporating causal relations in default reasoning schemes. Although inference rules implementing explaining away have been well studied [2], [9], precise and general

Manuscript received October 1991; revised September 1992. This work was supported by the National Science Foundation under grant IRI-8807061 to Carnegie Mellon and by the Rockwell International Science Center. Portions of this paper originally appeared in *Proceedings of the Second International Conference on Principles of Knowledge Representation and Reasoning* [16].

M. P. Wellman is with the Department of Electrical Engineering and Computer Science, University of Michigan, Ann Arbor, MI 48109.

M. Henrion is with Rockwell International Science Center, Palo Alto, CA, 94301.

IEEE Log Number 9206606.

¹By convention, upper-case letters denote propositional literals, whereas lower-case letters denote variables. Thus, variable *a*, “rain last night,” can take on the value *A* or its negation \bar{A} .




Fig. 1. Causal diagram for the “sprinkler” example [9].




Fig. 2. Drinking-and-driving example. Explaining away fails in this case because the two causes are positively related given their common effect.

conditions under which this pattern is valid have not appeared in the literature. Pearl provides these conditions for the special case of linear/Gaussian models (see p. 351 of [10]). Geffner provides a probabilistic justification of explaining away in terms of ϵ -semantics [1]. Both of these demonstrations are illustrative but do not capture the full range of situations in which such inference is appropriate.

Explaining away is an example of *intercausal inference* [3], that is, reasoning between two causes with a common effect, in contrast with pure causal or pure evidential reasoning. Although explaining away is often intuitively compelling, there are cases in which it appears inappropriate. Consider the following example, which is illustrated by the causal model of Fig. 2. You notice this newspaper headline about a well-known politician: “Senator Jones Killed in Car Accident.” You idly wonder whether she might have been drunk. The headline gives no indication of whether she was at fault in the accident or even whether she was a driver or passenger. You had no previous information about her driving or drinking habits, but you know that alcohol is a major cause of fatal car accidents. Reading on, you find out that Jones was indeed the driver and that no other vehicle was involved in the accident. How does this new information affect your belief that she had been drinking? Without knowledge of any accident, the fact that the Senator was driving might *reduce* the suspicion that she had been drinking. Given the accident, however, the fact that she was the driver would *increase* the suspicion. Note that this pattern of plausible reasoning is the *opposite* of explaining away: Knowledge of a common effect renders a positive dependence between the causes even though the causes were independent or even negatively dependent *a priori*.

The goal of this correspondence is to provide a general analysis of intercausal reasoning that accounts for both of the illustrated patterns of reasoning, and that makes precise the conditions differentiating them. Our choice of a probabilistic approach reflects the uncertainty that is central to causal explanation tasks and is supported by the observation that explaining away is a natural consequence of some generic structures commonly employed in probabilistic modeling.

Although the probabilistic formulation refers to quantitative degrees of belief, it does not necessarily require precise numerical probabilities for application. Indeed, our analysis is qualitative,

concerning the direction—but not the magnitude—of probabilistic dependencies. Our premise is that the critical distinctions correspond to intuitive categories of interaction among causes and that further precision would be impractical or less convenient and, for many purposes, unnecessary. This position is supported by the observation that common vocabulary includes numerous qualitative concepts of causal interaction. For example, we often say that causal factors act independently or synergistically, that one cause (a “gating condition”) enables or inhibits another, or that a set of available inputs are complementary or substitutable with one another. Rain and sprinkling independently cause wet grass; drinking amplifies the causal relation between driving and car accidents.

We formalize these concepts using the qualitative probabilistic network (QPN) representation [15], which is an abstraction of Bayesian networks. The analysis of intercausal reasoning extends this formalism by introducing new qualitative characterizations of causal interactions complementary with the existing QPN synergy relations.

In the remainder of this correspondence, we present a formal analysis of qualitative intercausal relations. After reviewing the notion of qualitative probabilistic influence in Section II, in Section III, we analyze intercausal reasoning with uncertain causal influences and identify the conditions under which explaining away will occur. We generalize these conditions in Section IV to handle prior intercausal relationships and partial evidence on the effect. Finally, in Section V, we present a view of Occam’s razor suggested by intercausal relations.

II. QUALITATIVE PROBABILISTIC NETWORKS

Our analysis of intercausal inference under uncertainty is based on the QPN formalism for qualitative probabilistic reasoning [15]. In a qualitative probabilistic network, variables are represented as nodes in a graph with directed edges defining probabilistic relationships. As in Bayesian networks [10] and other graphical schemes, connectedness in the graph represents the dependency structure of the underlying probability distribution [11]. However, rather than specify the distribution precisely with numeric probability tables, QPN’s merely constrain the conditional probabilities using qualitative influences. A sign $\delta \in \{+, -, 0, ?\}$ denoting the direction of qualitative influence between nodes is associated with each edge. Fig. 3 depicts an example QPN representing beliefs about the health of a friend. Event A , that our friend has a cold, increases² the probability of C , that he is sneezing. Event B , that he has an allergic reaction, also increases this probability. On the other hand, event F , that he recently took an antihistamine, reduces the probability of sneezing. Event D , that our friend is allergic to cats, increases the probability of an allergic reaction, as does E , that a cat is present. (Whereas, for ease of exposition, the variables in our examples are binary, the definition of qualitative influence that follows, like most other definitions and theorems, applies equally to multivalent discrete and continuous variables.)

For the general definition of qualitative influences, consider a QPN with a directed edge from a to c —and optionally some other variables collectively denoted x —with links to c . In Fig. 3, for example, x would comprise b and f . This structure dictates that the probability distribution for c can be specified conditionally on a and x .

Definition 1 (Qualitative Influence): We say that a positively influences c in a QPN G , which is written $S^+(a, c, G)$, if and only if (iff) for all values $a_1 > a_2$, c_0 , and all assignments x to other predecessors of c in G

$$\Pr(c \geq c_0 | a_1, x) \geq \Pr(c \geq c_0 | a_2, x).$$

²We use terms such as *increase* and *decrease* in the nonstrict sense, unless explicitly stated.




Fig. 3. Example QPN representing beliefs about a friend’s health. Arrows labeled “+” and “−” denote positive and negative causal influences, respectively.




Fig. 4. Schematic QPN transformation for intercausal inference. The qualitative influence δ of a on b on observation of C indicates whether explaining away occurs.

An equivalent condition is that the probability density function (or mass function in the discrete case) for a given c and x , $f_a(\cdot | c, x)$ obeys the *monotone likelihood-ratio property*:

$$\frac{f_a(a_1 | c_1, x)}{f_a(a_1 | c_2, x)} \geq \frac{f_a(a_2 | c_1, x)}{f_a(a_2 | c_2, x)} \quad (1)$$

for all $a_1 > a_2$, $c_1 > c_2$, and x . This property ensures that increasing a increases the expected value of c .³ We can replace \geq in (1) with \leq or $=$, yielding the conditions for negative influence S^- or zero influence S^0 , respectively. S^0 means that a and c are independent for all values of x . Note that the condition S^+ requires the inequality (1) to hold for all assignments x and similarly for S^- and S^0 . Therefore, it is quite possible that none of the three conditions hold. The condition $S^?$ indicates that the qualitative influence is ambiguous or that it is not known which, if any, of the relations holds.

III. PROBABILISTIC INTERCAUSAL RELATIONS

Suppose we observe our friend sneezing C , which raises the probability of his having a cold A , and the probability of his having an allergic reaction B . If we know that he is allergic to cats D , then learning that a cat is present E lends confirmation of the allergic reaction B . This explains away the sneezing and, therefore, reduces the probability of the cold A .

This process of intercausal inference can be cast as transformation of a causal graph or QPN, which is illustrated in general form in Fig. 4. Again, a and b are causes of c . For generality, we allow that there may be other causes of c (collectively represented by x) and that a and b in turn may have causal antecedents (b ’s are collectively labeled y ; a ’s do not figure in the example). Fig. 4(a) depicts this initial situation. Note that because their only connecting path is via direct links to c , a and b are marginally independent although they are conditionally dependent given c .

The basic explaining-away scenario starts with an observation of the effect variable to be explained c . Suppose that c is propositional

³In writing these ratios here and elsewhere, we assume that all conditional-probability terms are well defined and nonzero. These assumptions could be relaxed at the expense of explicatory complexity. For further discussion of these probabilistic inequalities, see [7] and [15].

and that the observed value is C . To represent observation in a probabilistic network, we instantiate the observed node and modify the dependency structure in the graph so that the nodes of interest become conditional on the observation. The evidence instantiation is tantamount to reversing the links from a and b to c [12], as shown in Fig. 4(b). The signs on the reversed links remain positive, indicating that observing C increases the probability of higher values of a and b . In addition, the reversals introduce a new *intercausal* link between a and b , accounting for the fact that the variables become dependent on observing C . The explaining-away pattern is characterized exactly by the negativity of this intercausal influence. For propositional a and b , the relation $S^-(a, b)$ in the graph of Fig. 4(b) would mean that

$$\Pr(B|ACxy) \leq \Pr(B|Cxy) \leq \Pr(B|\bar{A}Cxy).$$

Hence, belief in A decreases belief in B .

Even if we knew that the signs on the original links from a to c and b to c were positive, without further constraint, the sign on this new intercausal link would be ambiguous [15]. The question is this: What condition on the causal combination of a and b would enable us to derive a negative intercausal influence on observing C ?

Theorem 1 (Explaining Away): Let a and b be predecessors of c in a QPN G , and let x denote an assignment to c 's other predecessors, if any. Let $obs(c_0, G)$ denote the QPN obtained from G on observation of c_0 . Suppose $S^0(a, b, G)$. Then, $S^-(a, b, obs(c_0, G))$ iff for all $a_1 > a_2$, $b_1 > b_2$, and x

$$\frac{f_c(c_0|a_1b_1x)}{f_c(c_0|a_1b_2x)} \leq \frac{f_c(c_0|a_2b_1x)}{f_c(c_0|a_2b_2x)}. \quad (2)$$

This follows directly from Bayes's rule, reversing the dependence of c on b .⁴

Because it plays such a pivotal role in explaining away, we introduce terminology and notation for the intercausal relation (2).

Definition 2 (Product Synergy): Let a and b be predecessors of c in G , and let x denote an assignment to c 's other predecessors, if any. Variables a and b exhibit *negative product synergy* with respect to a particular value c_0 of c in G , which is written $X^-(\{a, b\}, c_0, G)$ if for all $a_1 > a_2$, $b_1 > b_2$ and x

$$f_c(c_0|a_1b_1x)f_c(c_0|a_2b_2x) \leq f_c(c_0|a_1b_2x)f_c(c_0|a_2b_1x). \quad (3)$$

Note that (3) is just the product form of (2). Thus, negative product synergy requires that the proportional increase in the probability of c_0 on raising b is smaller for higher values of a . Hence, the causal contribution of a given variable is greatest when that variable is the only active (high-valued) cause. It is this type of interaction that underlies explaining away.

We define *positive product synergy*, X^+ and *zero product synergy*, X^0 by substituting \geq and $=$, respectively, for \leq in (2). Theorem 1 is also valid with either "+" or "0" substituted for "-" in both the intercausal influence S^- and corresponding product synergy X^- . As for qualitative influences, the negative, zero, and positive product synergies are not exhaustive. The condition $X^?$ indicates that the product synergy is ambiguous or that it is not known which, if any, of the relations hold.

We illustrate the main result by reconsidering the two examples of explaining away. In Fig. 1, there is a negative intercausal relation between rain A and the sprinkler B , given their common effect, wet grass, C . Fig. 3 displays a corresponding negative intercausal relation between the cold and allergic reaction given their common effect, sneezing. According to Theorem 3, this kind of relationship is appropriate if and only if we believe that negative product synergy

⁴Complete proofs of this and other results are provided in the Appendix. A propositional version of Theorem 1 appears in [6].

holds in each of these cases, that is, our beliefs about the causal effects must satisfy

$$\frac{\Pr(C|AB)}{\Pr(C|\bar{A}B)} \leq \frac{\Pr(C|\bar{A}\bar{B})}{\Pr(C|A\bar{B})}. \quad (4)$$

In words, the proportional increase in probability of C , wet grass, due to learning B , sprinkling, is smaller given A , rain, than given \bar{A} , no rain. On the other hand, the proportional increase in the probability of sneezing due to learning that our friend has an allergy is less given a cold than given no cold. Both of these conditions seem eminently plausible—given one cause is present; the incremental effect of the second cause is less than it would be if the first were absent.

If negative product synergy does not seem immediately compelling, one can also derive it as a generalization of the *leaky noisy-OR* [4], [10], which is a plausible model for either situation. The noisy-OR dictates that each of the two causes may be sufficient alone to cause the effect and that the causal mechanisms are independent. The *leakiness* allows that even if neither A nor B occurs, C may occur for another unspecified reason (a *leak* L). It is easy to show that the leaky noisy-OR relation implies negative product synergy with respect to the presence of the effect and therefore leads to explaining away [16]. This result generalizes straightforwardly to cases with more than two causal variables. In contrast, *noisy-NOR* models—where causes lead to the *negation* of the effect—exhibit *zero* product synergy.

Now, let us reconsider examples for which explaining away does not seem to apply. The drinking and driving Senator from Fig. 2 is one such instance. The case from Fig. 3 of the two causes of an allergic reaction is another. Given that an allergic reaction B is observed, knowledge that our friend is allergic to cats D would tend to increase the probability that a cat is present E , and vice versa. There is a positive intercausal relationship between D and E , given B . According to the positive version of Theorem 1, this relationship holds iff positive product synergy applies—that is, iff

$$\frac{\Pr(B|DE)}{\Pr(B|\bar{D}\bar{E})} \geq \frac{\Pr(B|\bar{D}E)}{\Pr(B|D\bar{E})}. \quad (5)$$

For our example, this condition says that the proportional increase in probability of an allergic reaction due to the cat being present is greater, given that our friend is allergic to cats, than it would be if he were not. This is evident, given that the cat would have only indirect effects, if any, if he were not allergic to cats. Therefore, the right-hand side of (5) would be at or near unity, whereas the left-hand side would be significantly larger.

IV. EXTENSIONS

A. Dependent Causes

The premise of Theorem 1 requires that causes a and b be marginally independent. We can generalize the result, as long as any prior dependence between the causes is in the same direction as the intercausal effect of observing their common finding:

Theorem 2:

$$S^h(a, b, G) \wedge X^h(\{a, b\}, c_0, G) \Rightarrow S^h(a, b, obs(c_0, G)).$$

For example, suppose we know our neighbor habitually listens to weather reports and turns off the sprinkler when rain is forecast. This negative prior relation between the two causes is in the same direction as the intercausal relation, and hence, the tendency of the sprinkler to explain away the rain hypothesis is only strengthened.

On the other hand, suppose we believe in Murphy, the perverse rain god who likes to make it rain soon after a sprinkler has been used. This induces a positive prior dependence between the causes, rain and sprinkler. In this case, the intercausal relationship after observing




Fig. 5. Two causes a and b , with partial evidence e , for their common effect c .

wet grass becomes ambiguous and cannot be determined by purely qualitative analysis.

B. Indirect Evidence

Theorem 1 also presumes that the effect variable c is observed directly. Can we generalize the main result to situations where we have only indirect evidence for c ?

Suppose we observe the value of variable e , which is an effect of c . For example, in the sprinkler model, we might observe E , cold and shiny grass. To determine the intercausal implications of this observation, we investigate the interaction relation of a and b on e when c is factored out. This situation is depicted in Fig. 5.

To propagate intercausal reasoning through indirect evidence, we appeal to another synergy concept, which was previously introduced for QPN's [15]:

Definition 3 (Additive Synergy): Let a and b be predecessors of c in G , and let x denote an assignment to c 's other predecessors, if any. Variables a and b exhibit *negative additive synergy* with respect to variable c in G , which is written $Y^-(\{a, b\}, c, G)$, if for all $a_1 > a_2$, $b_1 > b_2$, x , and c_0

$$\Pr(c \geq c_0 | a_1 b_1 x) + \Pr(c \geq c_0 | a_2 b_2 x) \leq \Pr(c \geq c_0 | a_1 b_2 x) + \Pr(c \geq c_0 | a_2 b_1 x).$$

Positive additive synergy, Y^+ and *zero additive synergy* Y^0 are defined similarly, substituting \geq and $=$, respectively, for \leq . An important difference between additive and product synergy is that the former is defined with respect to the *variable* c , rather than to a particular value c_0 , that is, the additive synergy condition holds for all values of c . The disparity is due to the distinct roles of these relations in qualitative probabilistic inference. Note, however, that when c is a propositional variable, $Y^\delta(\{a, b\}, c)$ is identical to $X^\delta(\{a, b\}, C)$, except in substituting addition for multiplication in (3) (or differences for quotients in (2)). Although neither subsumes the other in general, when both of the individual influences of each cause on the effect have unambiguous signs (+ or -), then there are entailment relationships between them. See [16] for a detailed exposition of these relationships.

The following result establishes (for the propositional case) that evidence that is positively related to the effect maintains intercausal relations given some particular patterns of product and additive synergy.

Theorem 3: Let $red(c, G)$ denote the QPN obtained from G by reducing (averaging out) variable c . Suppose $X^{\delta_1}(\{a, b\}, C, G)$, $Y^{\delta_2}(\{a, b\}, c, G)$, $S^{\delta_3}(c, e, G)$, $S^0(a, e, G)$, and $S^0(b, e, G)$. Then, $X^{\delta_1}(\{a, b\}, E, red(c, G))$ if either of the following hold:

1. $\delta_1 = \delta_2$ and $\delta_3 = +$.
2. $\delta_1 = -\delta_2$ and $\delta_3 = -$.




Fig. 6. Multiple findings and complementary hypotheses: (a) QPN with three diseases that can cause two findings; (b) observing findings F^1 and F^2 . Explaining away produces two negative intercausal influences, which can be chained to reveal a positive relation between d^1 and d^3 .

Under certain circumstances, we can generalize Theorem 3 to the case of nonpropositional c . In essence, product synergy extends from c_0 to e_0 as long as e_0 supports c_0 but does not distinguish among $c \neq c_0$.⁵ For propositional c , it matters only whether the observed value e_0 was more likely given C than \bar{C} .

V. OCCAM'S RAZOR AND INTERCAUSAL REASONING

Suppose that there are several causal hypotheses—each of which could explain an observed effect by itself—related to the finding according to a negative product synergy relationship. Given the negative intercausal relations between each pair of hypotheses given the finding, invoking one hypothesis reduces belief in the others. This process is analogous to the action of Occam's razor in slicing away hypotheses that are multiplied beyond necessity.

On the other hand, if two or more causes interact with a *positive* product synergy, their joint occurrence may be a more likely explanation of the finding than would either of them alone. The synergistic effects of drinking and driving and of cat allergies and cats are two examples. We might be tempted to invoke "Occam's glue" in such cases as the multiple hypotheses adhere to each other to form a coherent scenario. Perhaps, however, it is more appealing to regard the conjunctive relation as suggesting their combination as a single compound hypothesis. Seen in this light, they are not being multiplied beyond necessity and are therefore not actually contravening the principle of parsimony.

Note that when there are multiple evidence variables, positive intercausal relationships and complementary hypotheses can arise even when all synergy relations are negative. Consider the QPN in Fig. 6(a), where three diseases—represented by propositional variables d^1 , d^2 , and d^3 —can variously account for two findings f^1 and f^2 . Suppose that all influences are positive and that the pairwise interactions satisfy negative product synergy. According to Theorem 1, given both findings F^1 and F^2 , we obtain the two negative intercausal influences $S^-(d^1, d^2)$ and $S^-(d^2, d^3)$, which is depicted in Fig. 6(b). Chaining these, we can conclude $S^+(d^1, d^3)$ on removal of d^2 , indicating that events D^1 and D^3 are complementary. This conclusion fits the intuitive observation that the findings can be explained either by the single disease D^2 or by the combination D^1 and D^3 . If D^1 and D^3 are common diseases and D^2 is relatively rare, it is quite possible that the combination is more probable than the single disease. Thus, the intercausal analysis dictates how causal events should be clustered in compound hypotheses. Events that are complementary in the causal explanation become related by positive influences without explicit set-covering computations.

Qualitative intercausal reasoning has also proven useful in the design of algorithms for quantitative probabilistic diagnosis. Because exact inference is intractable for large multiply-connected networks, there has been considerable interest in approximation algorithms. One

⁵To establish this, we divide c into values for which X^δ holds (C) and those for which it does not (\bar{C}) and then apply the previous theorem. In the process, we must be careful that the division does not invalidate the conditional independence of a and b from e given c .

such approach for diagnosis is to use heuristic search to find the most probable hypotheses that can explain the observed findings. In one large medical diagnosis application—quick medical reference-belief network (QMR-BN) [13]—there are almost 600 diseases and, hence, 2^{600} potential diagnoses. However, in most cases, only a fraction of these diagnoses have substantial probability. Search-based algorithms, such as TopN [5], concentrate on the most probable hypotheses. Given the relative probabilities of the candidate diagnoses, TopN computes bounds on their absolute probabilities. The bounds may be successively narrowed as the search continues.

The key to the design of efficient search-based algorithms is an *admissibility heuristic* that allows them to prune subtrees that can provably lead only to hypotheses whose probability is less than some threshold. The TopN algorithm starts out by examining single-disease hypotheses and extending them incrementally. Intercausal analysis can identify which additional diseases are complementary and can therefore possibly lead to more probable hypotheses. It also reveals which diseases are competitive and can therefore lead only to less probable hypotheses. Thus, intercausal analysis provides a suitable basis for an admissibility heuristic. Because QMR-BN uniformly assumes noisy-OR relations among diseases and findings, the diseases are always competitive. Initial results for this network using this pruning criterion show rapid convergence to narrow probability bounds in most cases [5]. The analysis described in this paper generalizes this approach to handle networks not only with noisy-OR relations, as in QMR-BN, but with any interactions satisfying negative product synergy.

VI. CONCLUSIONS

Intercausal relations play a central role in the combination of diagnostic and predictive reasoning. The qualitatively significant property of interacting hypotheses is whether they compete with or complement one another in explaining the observed findings. In the former case, one cause explains away the other, given the observation. In addition, we have shown that explaining away is not the only pattern of intercausal reasoning. To account for this distinction, we have derived a general probabilistic criterion (*negative product synergy*) that precisely justifies explaining away.

The main appeal of qualitative probabilistic relations is that they require minimal precision yet capture some of the most significant behaviors. However, qualitative probabilistic inference may be useful, even for numerical systems, as a means of explanation to human users in a way that might correspond more directly to intuitive categories [6].

We also believe that it may be computationally advantageous to maintain these qualitative distinctions even when numeric information is available. As described in Section V, intercausal relations qualitatively restrict the reasonable patterns in which to cluster events in compound hypotheses. These constraints can be exploited in diagnosis to prune the space of composite hypotheses at a high level, based on qualitative admissibility.

APPENDIX PROOFS

Theorem 1: Let a and b be predecessors of c in a QPN G . Let $obs(c_0, G)$ denote the QPN obtained from G on observation of $c = c_0$. Suppose $S^0(a, b, G)$. Then, $S^-(a, b, obs(c_0, G))$ iff $X^-(\{a, b\}, c_0, G)$.

Proof: Let y denote the predecessors of b in G and x the predecessors of c other than a and b , if any. The distribution for

a given b, x , and y on observation of c_0 is, by Bayes's rule

$$f_a(a|bc_0xy) = \frac{f_c(c_0|abxy)f_a(a|bxy)}{f_c(c_0|bxy)}. \quad (6)$$

By conditional independence, we can drop the y condition from the f_c terms and the x condition from the f_a term on the right-hand side. The qualitative influence of a on b is positive iff (6) obeys the monotone likelihood ratio property (1) and negative iff the inequality of (1) is reversed. Substituting (6) in the likelihood ratio for a_i given a pair of values for $b, b_1 > b_2$, we obtain

$$\frac{f_c(c_0|a_i b_1 x)f_a(a_i|b_1 y)f_c(c_0|b_2 x)}{f_c(c_0|a_i b_2 x)f_a(a_i|b_2 y)f_c(c_0|b_1 x)}. \quad (7)$$

Since $f_c(c_0|b_j x)$ does not depend on a_i , the ratio (7) is increasing or decreasing in a_i in direct correspondence with

$$\frac{f_c(c_0|a_i b_1 x)f_a(a_i|b_1 y)}{f_c(c_0|a_i b_2 x)f_a(a_i|b_2 y)}. \quad (8)$$

By the conditional independence of a and b given y (the S^0 condition), $f_a(a_i|b_1 y) = f_a(a_i|b_2 y)$; therefore, these terms may be canceled from the expression, leaving

$$\frac{f_c(c_0|a_i b_1 x)}{f_c(c_0|a_i b_2 x)}.$$

The direction of change of this expression with respect to a_i is exactly the product synergy condition. \square

Theorem 2:

$$S^\delta(a, b, G) \wedge X^\delta(\{a, b\}, c_0, G) \Rightarrow S^\delta(a, b, obs(c_0, G)).$$

Proof: Proceed as for Theorem 1 up to the reference to unconditional independence. The ratio (8) can be factored into two parts:

$$\left(\frac{f_c(c_0|a_i b_1 x)}{f_c(c_0|a_i b_2 x)} \right) \left(\frac{f_a(a_i|b_1 y)}{f_a(a_i|b_2 y)} \right).$$

The first part increases according to the sign of product synergy, and the second is contingent on the direct influence of a on b prior to observation of c_0 . When the two agree, the direction of the entire expression is determined, establishing the qualitative influence of a on b posterior to the observation. \square

Theorem 3: Let $red(c, G)$ denote the QPN obtained from G by reducing (averaging out) variable c . Suppose $X^{\delta_1}(\{a, b\}, C, G)$, $X^{\delta_2}(\{a, b\}, c, G)$, $S^{\delta_3}(c, e, G)$, $S^0(a, e, G)$, and $S^0(b, e, G)$. Then, $X^{\delta_1}(\{a, b\}, E, red(c, G))$ if either of the following occur:

1. $\delta_1 = \delta_2$ and $\delta_3 = +$.
2. $\delta_1 = -\delta_2$ and $\delta_3 = -$.

Proof: Let $H_{i,j} = \Pr(E|a_i b_j x)$ and $G_{i,j} = \Pr(C|a_i b_j x)$. Since e is conditionally independent of a and b given c

$$H_{i,j} = \Pr(E|C)G_{i,j} + \Pr(E|\bar{C})(1 - G_{i,j}).$$

Expanding terms and simplifying, the product of two H expressions is

$$H_{i,j}H_{k,l} = G_{i,j}G_{k,l}\Delta^2 + (G_{i,j} + G_{k,l})\Pr(E|\bar{C})\Delta$$

where $\Delta = [\Pr(E|C) - \Pr(E|\bar{C})]$, which is positive or negative according to δ_3 . Since Δ^2 is always positive, the comparison of a pair of H products is the same as for the corresponding G products if the comparison of second additive terms also agrees. When $\Delta > 0$, the sign of this second comparison is determined by the additive synergy relation and when $\Delta < 0$ by its negation. \square

ACKNOWLEDGMENT

We thank M. Druzdzel for discussions on several aspects of this work and L. Dupre for technical editing. R. Patil and the anonymous reviewers provided useful suggestions on the presentation of this material.

REFERENCES

- [1] H. Geffner, "On the logic of defaults," in *Proc. Nat. Conf. Artificial Intell.* (St. Paul, MN), 1988, pp. 449-454.
- [2] —, "Causal theories for nonmonotonic reasoning," in *Proc. Nat. Conf. Artificial Intell.* (Boston, MA), 1990, pp. 524-530.
- [3] M. Henrion, "Uncertainty in artificial intelligence: Is probability epistemologically and heuristically adequate?" in *Expert Judgment and Expert Systems (NATO ISI Series F)* (J. Mumpower et al., Eds.). Berlin: Springer-Verlag, 1987, pp. 105-130, vol. 35.
- [4] —, "Some practical issues in constructing belief networks," in *Uncertainty in Artificial Intelligence 3* (L. N. Kanal, T. S. Levitt, and J. F. Lemmer, Eds.). Amsterdam: North-Holland, 1989.
- [5] —, "Search-based methods to bound diagnostic probabilities in very large belief nets," in *Proc. Seventh Conf. Uncertainty Artificial Intell.* (Los Angeles, CA.), 1991, pp. 142-150.
- [6] M. Henrion and M. J. Druzdzel, "Qualitative propagation and scenario-based explanation of probabilistic reasoning," in *Uncertainty in Artificial Intelligence 6* (P. P. Bonissone, M. Henrion, and L. N. Kanal, Eds.). Amsterdam: North-Holland, 1991.
- [7] P. R. Milgrom, "Good news and bad news: Representation theorems and applications," *Bell J. Econ.*, vol. 12, pp. 380-391, 1981.
- [8] E. Paek, "A circumscriptive theory for causal and evidential support," in *Proc. Nat. Conf. Artificial Intell.*, 1990, pp. 545-549.
- [9] J. Pearl, "Embracing causality in default reasoning," *Artificial Intell.*, vol. 35, pp. 259-271, 1988.
- [10] —, *Probabilistic Reasoning in Intelligent Systems: Networks of Plausible Inference*. San Mateo, CA: Morgan Kaufmann, 1988.
- [11] J. Pearl, D. Geiger, and T. Verma, "Conditional independence and its representations," *Kybernetika*, vol. 25, pp. 33-44, 1989.
- [12] R. D. Shachter, "Evidence absorption and propagation through evidence reversals," in *Proc. Workshop Uncertainty Artificial Intell.* (Windsor, Canada), 1989, pp. 303-310.
- [13] M. Shwe, B. Middleton, and D. E. Heckerman, "Probabilistic diagnosis using a reformulation of the Internist-I/QMR knowledge base: I. The probabilistic model and inference algorithms," *Methods Inform. Med.*, vol. 30, pp. 241-255, 1991.
- [14] M. P. Wellman, *Formulation of Tradeoffs in Planning Under Uncertainty*. London: Pitman, 1990.
- [15] —, "Fundamental concepts of qualitative probabilistic networks," *Artificial Intell.*, vol. 44, pp. 257-303, 1990.
- [16] M. P. Wellman and M. Henrion, "Qualitative intercausal relations, or explaining 'explaining away,'" in *Principles Knowledge Represent. Reasoning: Proc. Sec. Int. Conf.*, 1991, pp. 535-546.

An Approximate Nonmyopic Computation for Value of Information

David Heckerman, Eric Horvitz, and Blackford Middleton

Abstract—Value-of-information analyses provide a means for selecting the next best observation to make and for determining whether it is better to gather additional information or to act immediately. Determining the next best test to perform, given uncertainty about the state of the world, requires a consideration of the value of making all possible sequences of observations. In practice, decision analysts and expert-system designers have avoided the intractability of exact computation of the value of information by relying on a *myopic* assumption that only one additional test will be performed, even when there is an opportunity to make a large number of observations. We present an alternative to the myopic analysis. In particular, we present an approximate method for computing the value of information of a set of tests, which exploits the statistical properties of large samples. The approximation is linear in the number of tests, in contrast with the exact computation, which is exponential in the number of tests. The approach is not as general as is a complete nonmyopic analysis, in which all possible sequences of observations are considered. In addition, the approximation is limited to specific classes of dependencies among evidence and to binary hypothesis and decision variables. Nonetheless, as we demonstrate with a simple application, the approach can offer an improvement over the myopic analysis.

Index Terms—Belief networks, decision theory, nonmyopic, probability, value of information.

I. INTRODUCTION

When performing diagnosis, a person usually has the opportunity to gather additional information about the state of the world before making a final diagnosis. Such information gathering typically is associated with costs and benefits. These costs and benefits can be balanced with decision-theoretic techniques—in particular, procedures for computing *value of information*. These techniques form an integral part of many decision-theoretic expert systems for diagnosis, such as Gorry and Barnett's program for the diagnosis of congestive heart failure [1].

In most diagnosis contexts, a decisionmaker has the option to perform several tests and can decide which test to perform after seeing the results of all previous tests. Thus, a person or expert system should consider the value of all possible sequences of tests. Such an analysis is intractable because the number of sequences grows exponentially with the number of tests. Builders of expert systems have avoided the intractability of exact value-of-information computations by implementing *myopic* or *greedy* value-of-information analyses. In such analyses, a system determines the next best test by computing the value of information based on the assumption that the decisionmaker will act immediately after seeing the results of the single test [2].

Manuscript received November 1991; revised August 1992. This work was supported by the National Cancer Institute under Grant RO1CA51729-01A1 and by the Agency for Health Care Policy and Research under Grant T2HS00028.

D. Heckerman is with the Department of Computer Science, University of California, Los Angeles, CA 90024-1596. He is also with Microsoft Research Center, Redmond, WA 98052-6399.

E. Horvitz is with the Palo Alto Laboratory, Rockwell International Science Center, Palo Alto, CA 94301. He is also with the Section on Medical Informatics, Stanford University Medical Center, Stanford, CA 94305-5479.

B. Middleton is with Clinical Information Systems, Stanford University Medical Center, Stanford, CA 94305-5479.

IEEE Log Number 9206605.